Two Day Meeting Preparation Plan

Prepare 18-19 October | Meet 20 October 2026 | All times Asia/Shanghai

Meeting goal: leave with one agreed handoff between modules, working evidence from each person, a clear owner for every blocker, and the first integrated balance flow. The whole product does not need to be finished in two days.

This is a separate short-term action plan. The exact submission deadline and meeting time still need confirmation. The earlier 12-day roadmap remains the broader direction.

What the current checkout tells us

P1 has a tested standalone agent. P2 has a backend still using mocks, but no frontend implementation in this checkout. P3 has working analytics. P5 has cancellation eligibility but no merged detector. P4 has no merged implementation here; ask them to share their current branch before assigning duplicate work.

The reviewed combined test suite passed 148 tests. The tests do not prove that the real modules work together. The highest priority is making their inputs, outputs and IDs agree.

Required outcomes by the meeting

Person	Minimum to bring	Main teammate handoff
P1 Agent	Clarification demo, canonical proposals, proposed context/result examples	P2: conversation state and tool results
P2 App - Marvel	Backend alignment plus one balance flow using P1; documented startup	P1: adapter; all: IDs, API and seed data
P3 Finance	Reproducible reports, exact expected totals, data-history requirements	P2: query input; P1: report explanation
P4 Security	Reviewable policy branch, decision examples and passing policy tests	P2: enforcement; P1: user-facing outcomes
P5 Subscriptions	Minimal detector with evidence, candidates and eligibility examples	P2: ID mapping; P1: candidate explanations

Shared rules for the next 48 hours

Today: Each person shares their branch name, latest commit, start/test commands, and one blocking question. Share draft handoff JSON now; do not wait until the meeting.

Before the meeting: Push a reviewable commit/PR. Run your own tests. Bring a short terminal/API recording or saved output if a live demo is unreliable. Never include API keys or real customer data.

Ownership: Keep edits in assigned directories. A shared-contract change needs maintainer approval and versioning. If unresolved, bring a proposed example; do not silently redefine shared DTOs.

Definition of ready: A teammate can run your example from its instructions. Say clearly whether each result is real module output, stubbed, or still blocked. No polished slides are needed yet.

P1 and P2 before the meeting

Build one reliable handoff together; avoid two competing backends

P1 Agent - you

18 October | Must finish: Send P2 one canonical transfer proposal and one balance proposal using the frozen schemas. Include exact catalog/context and tool-result examples from docs/agent/integration.md. Mark which conventions still need agreement.

18-19 October | Must finish: Build a small local conversation demonstration: transfer to Alice, full-name clarification, then amount clarification, then one valid proposal. Keep state in the harness/caller rather than a global service variable. Preserve details already supplied.

19 October | Must finish: Show that never mind clears a pending request and that two separate conversation IDs do not share state. Add regression tests for those exact cases. If either is unfinished, bring its failing transcript and a bounded fix plan.

With Marvel: Help connect the same public process_message interface to one backend balance lookup. A mock provider is acceptable for the first integration; repeat with Gemini if accessible. Bank data must come from the backend, not Gemini.

Bring: One successful conversation transcript; one interrupted conversation; proposal/context/result examples; passing tests or named failures; branch/commit and exact run command.

Stretch only: Better report wording and additional paraphrases. Do not build SQL persistence, a full workflow engine, or a second app.

P2 App and backend - Marvel

18 October | Must finish: Align the backend transfer input with recipient_id; remove the fallback Bob and use the proposed source account only after ownership checks. Use shared argument validation before risk evaluation. Check the seeded IDs with P1/P3/P5.

18 October | Must finish: Implement a real/mock agent switch; merely changing use_mock currently does not switch the call. Accept an authenticated catalog/context and document how one correlated result returns to P1.

19 October | Must finish: Make balance work through POST /chat: P1 proposes get_balance; P2 reads the selected account; P2 returns a typed result; P1 explains it. Read-only lookup must not create a transfer approval or call execute_transfer.

19 October | Agree with team: Show the proposed approval response shape and API wrappers. Keep action_id/digest/tier separate from strict tool arguments through an agreed schema. Bring any necessary versioned contract change to the meeting.

Bring: Backend startup and safe synthetic seed/reset instructions; one working balance request; current transfer-confirmation/MFA demonstration; API examples and tests. Explain that the demo seed reset drops/recreates demo tables.

Stretch only: A bare chat/balance screen with loading/error states. If time is tight, FastAPI docs are enough for the meeting; the frontend must begin immediately afterward.

Meeting acceptance: P1 and P2 can trace the same request from message to proposal to database result to final reply, with matching conversation/account IDs. If the connection is blocked, identify the exact differing field and agree its owner before adding features.

P3 P4 and P5 before the meeting

Bring concrete domain outputs and integration decisions

P3 Finance analytics

18 October | Must finish: Prepare a short input checklist for P2: selected-account transactions, requested month plus three preceding months for anomalies, validated Decimal values, and Asia/Shanghai boundaries. State explicitly that the function aggregates every account passed to it.

19 October | Must finish: Produce two saved reports: the shared October fixture with total CNY 1843.00, and an anomaly-bearing example from your synthetic data. Include expected category totals and the transactions explaining each anomaly.

With P1 + P2: Confirm gross-spending rules: credits are excluded, refunds do not reduce spending, outbound transfers count by default. Keep general diagnostics outside SpendingReport for now; decide who displays them without inventing a notes field.

Bring: Both report JSON examples, expected totals, history/timezone instructions, test command/results, and one explanation of CODE: reason formatting. No new analytics features are required before this meeting.

P4 Security and risk

18 October | Must finish: Share the current branch/commit and a reviewable policy module. Implement or confirm the frozen evaluate_action(proposal, context) signature. Start with pure fail-closed rules, not backend/database code.

19 October | Must finish: Show decisions for read-only Green, CNY 500 with daily total 0 Yellow, CNY 500 with daily total 800 Red, exactly CNY 1000 cumulative, locked-account mutation, and invalid arguments. Distinguish confirmation from MFA and final authorization.

With P2 + P1: Agree which checks must be repeated at execution, action expiry and failed-MFA behavior, and user-facing reason codes. Identify which protections belong to policy versus P2 persistence/execution.

Bring: Decision JSON for those cases, policy/adversarial tests toward the original 20-case target, a short known-risk list, and branch/PR. A backend stub cannot count as the real P4 module.

P5 Subscription detection and eligibility

18 October | Must finish: Provide the frozen detect_subscriptions(transactions) -> SubscriptionReport entry point. Find clear repeated merchant/amount/cadence patterns and return confidence plus evidence_transaction_ids. Keep calculation separate from execution.

19 October | Must finish: Demonstrate a clear recurring candidate, weak/one-off evidence that is not asserted active, and unsupported cancellation. Cover duplicate/noisy inputs and deterministic output with tests toward the original 12-case target.

With P2 + P1: Agree how candidate IDs map to stored subscriptions and who supplies cancellation_supported/active status. Never invent a cancellable subscription record just because payments repeat. Current eligibility checks only the support flag; assign the remaining checks explicitly.

Bring: Candidate report JSON with evidence, supported/unsupported eligibility examples, ID-mapping proposal, test results, and branch/PR. No real merchant cancellation API is needed.

Meeting agenda and work afterward

20 October | Suggested 90-minute working session

Minutes	Activity	Required outcome
0-10	All: confirm deadline and current commits	Agree final submission date, demo scope and remaining working days
10-25	P1 + P2: trace one balance conversation	One working handoff or a precise blocker with owner
25-40	P3 + P2 + P1: report demonstration	Agree account/history/timezone input and displayed report fields
40-55	P5 + P2 + P1: candidate demonstration	Agree evidence, stored ID mapping and supported cancellation
55-70	P4 + P2: policy and approval demonstration	Agree validation order, thresholds and execution checks
70-85	All: freeze interface examples and priorities	Record approved schemas, decisions, owners and next acceptance tests
85-90	All: confirm next integration checkpoint	Choose a check-in within 48 hours and list what must run then

Decisions to write down before leaving

1. One synthetic dataset and ID mapping. 2. Catalog/context/result shape and conversation ownership. 3. Shared API/approval schemas and maintainer for any version bump. 4. Account scope, history window and reporting timezone. 5. Candidate versus active subscription rules. 6. Exact confirmation/MFA boundaries and safe failure behavior. 7. Demo scope, final deadline and first feature-freeze date.

After the meeting

Next 2 days | First integrated scenarios: P2 connects real P1/P3/P4/P5 calls and starts the minimal UI. P1 finishes the agreed conversation loop and grounded report replies. P3/P4/P5 fix their own integration mismatches and verify expected results. Target: balance and spending work through the app; transfer and subscription handoffs are exercised.

Following 2 days | Three scenarios and one connected sequence: Finish transfer confirmation/MFA and supported simulated cancellation. P1/P2 then connect detection, cancellation completion and a separately approved transfer; P4 checks every mutation, P5 supplies evidence, P3 checks projected-savings wording. Do not treat future savings as a credited balance.

Following 2 days | Hardening and deployment: P4 leads adversarial checks; P2 fixes execution-time ownership/lock/funds/policy/expiry, concurrent replay protection and audit, and tests the deployment database. All owners test failure recovery and a reproducible seed/reset. Freeze new features.

Final 2 days before submission | Deliver and rehearse: All prepare the brief (up to 200 words), technical/security documentation, source/README/tests, presentation (up to 10 minutes) and video (up to 5 minutes). Rehearse a normal flow, approval/MFA and a safe failure. Keep a clearly labelled mock demo available.

If behind: complete three reliable core scenarios before extras. Reduce connected-workflow breadth and visual polish; do not remove validation, consent, truthful results or ownership checks. If the actual deadline is earlier, shorten intermediate work and protect submission time. At every check-in, name one owner and next action for each blocker.